

1_Data_Exploration

April 25, 2025

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import sys
import os
```

```
[2]: # --- Configuration ---
# Path to the raw data file
RAW_DATA_PATH = "../data/raw/YearPredictionMSD.txt"
# Directory to save generated plots
PLOT_SAVE_DIR = "../results/plots/"
# Ensure plot directory exists
os.makedirs(PLOT_SAVE_DIR, exist_ok=True)

# --- Optional: Add src to path to reuse functions ---
module_path = os.path.abspath(os.path.join('.', 'src'))
if module_path not in sys.path:
    sys.path.append(module_path)
from data_processing import load_data, create_decade_bins
#If not reusing functions from src, redefine them or necessary parts below
```

```
[3]: # --- Data Loading ---
print(f"Loading raw data from: {RAW_DATA_PATH}")
try:
    # Define column names as in data_processing.py
    N_FEATURES = 90
    colnames = ['Year'] + [f'Feature_{i+1}' for i in range(N_FEATURES)]
    df_raw = pd.read_csv(RAW_DATA_PATH, header=None, names=colnames)
    print(f"Data loaded successfully. Shape: {df_raw.shape}")
    print("\nFirst 5 rows of raw data:")
    print(df_raw.head())
    print("\nBasic data info:")
    df_raw.info()
except FileNotFoundError:
    print(f"ERROR: Raw data file not found at {RAW_DATA_PATH}. Please ensure_
↳ it's downloaded.")
    # Exit or handle error appropriately in a real script
```

```

# For a notebook, we might just stop execution here or raise the error
raise
except Exception as e:
    print(f"An error occurred during data loading: {e}")
    raise

```

Loading raw data from: ../data/raw/YearPredictionMSD.txt
Data loaded successfully. Shape: (515345, 91)

First 5 rows of raw data:

	Year	Feature_1	Feature_2	Feature_3	Feature_4	Feature_5	Feature_6	\
0	2001	49.94357	21.47114	73.07750	8.74861	-17.40628	-13.09905	
1	2001	48.73215	18.42930	70.32679	12.94636	-10.32437	-24.83777	
2	2001	50.95714	31.85602	55.81851	13.41693	-6.57898	-18.54940	
3	2001	48.24750	-1.89837	36.29772	2.58776	0.97170	-26.21683	
4	2001	50.97020	42.20998	67.09964	8.46791	-15.85279	-16.81409	

	Feature_7	Feature_8	Feature_9	...	Feature_81	Feature_82	Feature_83	\
0	-25.01202	-12.23257	7.83089	...	13.01620	-54.40548	58.99367	
1	8.76630	-0.92019	18.76548	...	5.66812	-19.68073	33.04964	
2	-3.27872	-2.35035	16.07017	...	3.03800	26.05866	-50.92779	
3	5.05097	-10.34124	3.55005	...	34.57337	-171.70734	-16.96705	
4	-12.48207	-9.37636	12.63699	...	9.92661	-55.95724	64.92712	

	Feature_84	Feature_85	Feature_86	Feature_87	Feature_88	Feature_89	\
0	15.37344	1.11144	-23.08793	68.40795	-1.82223	-27.46348	
1	42.87836	-9.90378	-32.22788	70.49388	12.04941	58.43453	
2	10.93792	-0.07568	43.20130	-115.00698	-0.05859	39.67068	
3	-46.67617	-12.51516	82.58061	-72.08993	9.90558	199.62971	
4	-17.72522	-1.49237	-7.50035	51.76631	7.88713	55.66926	

	Feature_90
0	2.26327
1	26.92061
2	-0.66345
3	18.85382
4	28.74903

[5 rows x 91 columns]

Basic data info:

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 515345 entries, 0 to 515344
Data columns (total 91 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Year        515345 non-null int64
1   Feature_1   515345 non-null float64

```

2	Feature_2	515345	non-null	float64
3	Feature_3	515345	non-null	float64
4	Feature_4	515345	non-null	float64
5	Feature_5	515345	non-null	float64
6	Feature_6	515345	non-null	float64
7	Feature_7	515345	non-null	float64
8	Feature_8	515345	non-null	float64
9	Feature_9	515345	non-null	float64
10	Feature_10	515345	non-null	float64
11	Feature_11	515345	non-null	float64
12	Feature_12	515345	non-null	float64
13	Feature_13	515345	non-null	float64
14	Feature_14	515345	non-null	float64
15	Feature_15	515345	non-null	float64
16	Feature_16	515345	non-null	float64
17	Feature_17	515345	non-null	float64
18	Feature_18	515345	non-null	float64
19	Feature_19	515345	non-null	float64
20	Feature_20	515345	non-null	float64
21	Feature_21	515345	non-null	float64
22	Feature_22	515345	non-null	float64
23	Feature_23	515345	non-null	float64
24	Feature_24	515345	non-null	float64
25	Feature_25	515345	non-null	float64
26	Feature_26	515345	non-null	float64
27	Feature_27	515345	non-null	float64
28	Feature_28	515345	non-null	float64
29	Feature_29	515345	non-null	float64
30	Feature_30	515345	non-null	float64
31	Feature_31	515345	non-null	float64
32	Feature_32	515345	non-null	float64
33	Feature_33	515345	non-null	float64
34	Feature_34	515345	non-null	float64
35	Feature_35	515345	non-null	float64
36	Feature_36	515345	non-null	float64
37	Feature_37	515345	non-null	float64
38	Feature_38	515345	non-null	float64
39	Feature_39	515345	non-null	float64
40	Feature_40	515345	non-null	float64
41	Feature_41	515345	non-null	float64
42	Feature_42	515345	non-null	float64
43	Feature_43	515345	non-null	float64
44	Feature_44	515345	non-null	float64
45	Feature_45	515345	non-null	float64
46	Feature_46	515345	non-null	float64
47	Feature_47	515345	non-null	float64
48	Feature_48	515345	non-null	float64
49	Feature_49	515345	non-null	float64

```

50 Feature_50 515345 non-null float64
51 Feature_51 515345 non-null float64
52 Feature_52 515345 non-null float64
53 Feature_53 515345 non-null float64
54 Feature_54 515345 non-null float64
55 Feature_55 515345 non-null float64
56 Feature_56 515345 non-null float64
57 Feature_57 515345 non-null float64
58 Feature_58 515345 non-null float64
59 Feature_59 515345 non-null float64
60 Feature_60 515345 non-null float64
61 Feature_61 515345 non-null float64
62 Feature_62 515345 non-null float64
63 Feature_63 515345 non-null float64
64 Feature_64 515345 non-null float64
65 Feature_65 515345 non-null float64
66 Feature_66 515345 non-null float64
67 Feature_67 515345 non-null float64
68 Feature_68 515345 non-null float64
69 Feature_69 515345 non-null float64
70 Feature_70 515345 non-null float64
71 Feature_71 515345 non-null float64
72 Feature_72 515345 non-null float64
73 Feature_73 515345 non-null float64
74 Feature_74 515345 non-null float64
75 Feature_75 515345 non-null float64
76 Feature_76 515345 non-null float64
77 Feature_77 515345 non-null float64
78 Feature_78 515345 non-null float64
79 Feature_79 515345 non-null float64
80 Feature_80 515345 non-null float64
81 Feature_81 515345 non-null float64
82 Feature_82 515345 non-null float64
83 Feature_83 515345 non-null float64
84 Feature_84 515345 non-null float64
85 Feature_85 515345 non-null float64
86 Feature_86 515345 non-null float64
87 Feature_87 515345 non-null float64
88 Feature_88 515345 non-null float64
89 Feature_89 515345 non-null float64
90 Feature_90 515345 non-null float64
dtypes: float64(90), int64(1)
memory usage: 357.8 MB

```

```

[4]: # --- Decade Binning (Reproduce logic from data_processing.py) ---
print("\nCreating decade bins for analysis...")
min_year = 1920 # Start decade reference

```

```

df_raw['Decade_Start'] = (df_raw['Year'] // 10) * 10
df_raw['Decade_Label'] = ((df_raw['Decade_Start'] - min_year) // 10).astype(int)
df_raw['Decade_Label'] = df_raw['Decade_Label'].clip(lower=0) # Clip years < 1920
decade_map = {i: f"{min_year + i*10}s" for i in range(10)}
df_raw['Decade_Name'] = df_raw['Decade_Label'].map(decade_map)
print("Decade columns ('Decade_Label', 'Decade_Name') added.")

```

Creating decade bins for analysis...

Decade columns ('Decade_Label', 'Decade_Name') added.

```

[5]: # --- 1. Class Balance Analysis ---
print("\n--- 1. Class Balance Analysis ---")
plt.figure(figsize=(10, 6))
sns.countplot(data=df_raw, x='Decade_Name', order=[decade_map[i] for i in range(10)], palette='viridis')
plt.title('Distribution of Songs Across Decades')
plt.xlabel('Decade')
plt.ylabel('Number of Songs')
plt.xticks(rotation=45)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.savefig(os.path.join(PLOT_SAVE_DIR, 'eda_decade_distribution.png'))
plt.show()

decade_counts = df_raw['Decade_Name'].value_counts().sort_index()
print("\nSong Counts per Decade:")
print(decade_counts)
print(f"\nObservations: The dataset is heavily imbalanced, with a vast majority of songs from the 2000s, followed by the 1990s. Earlier decades have significantly fewer samples.")

```

--- 1. Class Balance Analysis ---

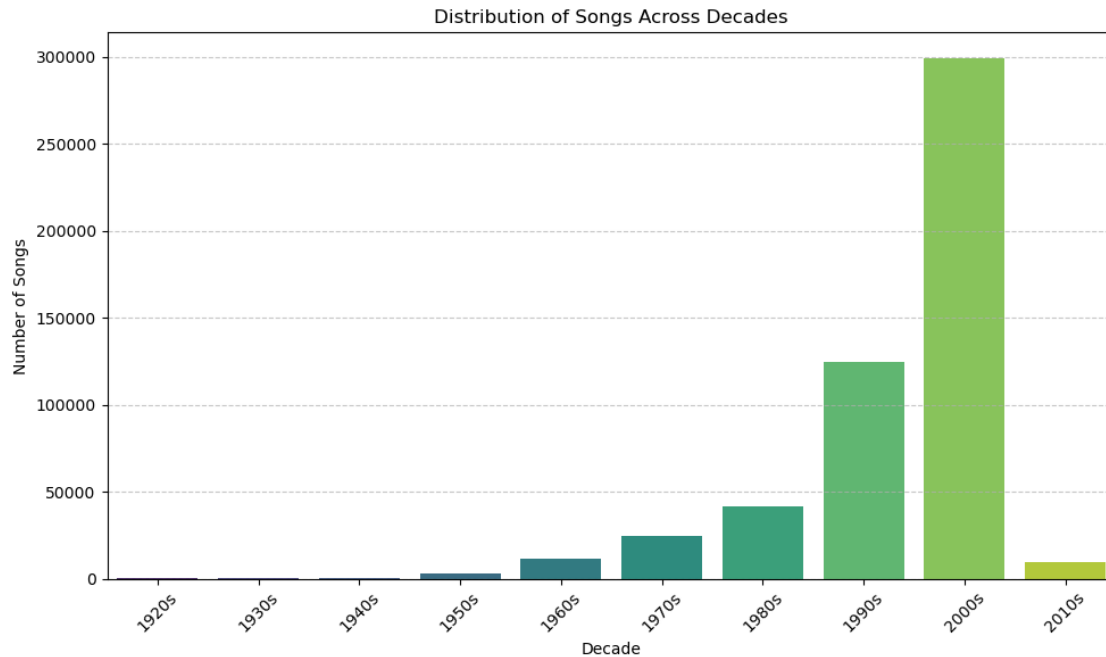
/var/folders/yf/9gf2xg3j2q76vf7cw8s4v4tw0000gn/T/ipykernel_10423/3132200938.py:4
: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

sns.countplot(data=df_raw, x='Decade_Name', order=[decade_map[i] for i in range(10)], palette='viridis')

```



Song Counts per Decade:

Decade_Name

1920s	224
1930s	252
1940s	356
1950s	3102
1960s	11739
1970s	24745
1980s	41814
1990s	124713
2000s	299003
2010s	9397

Name: count, dtype: int64

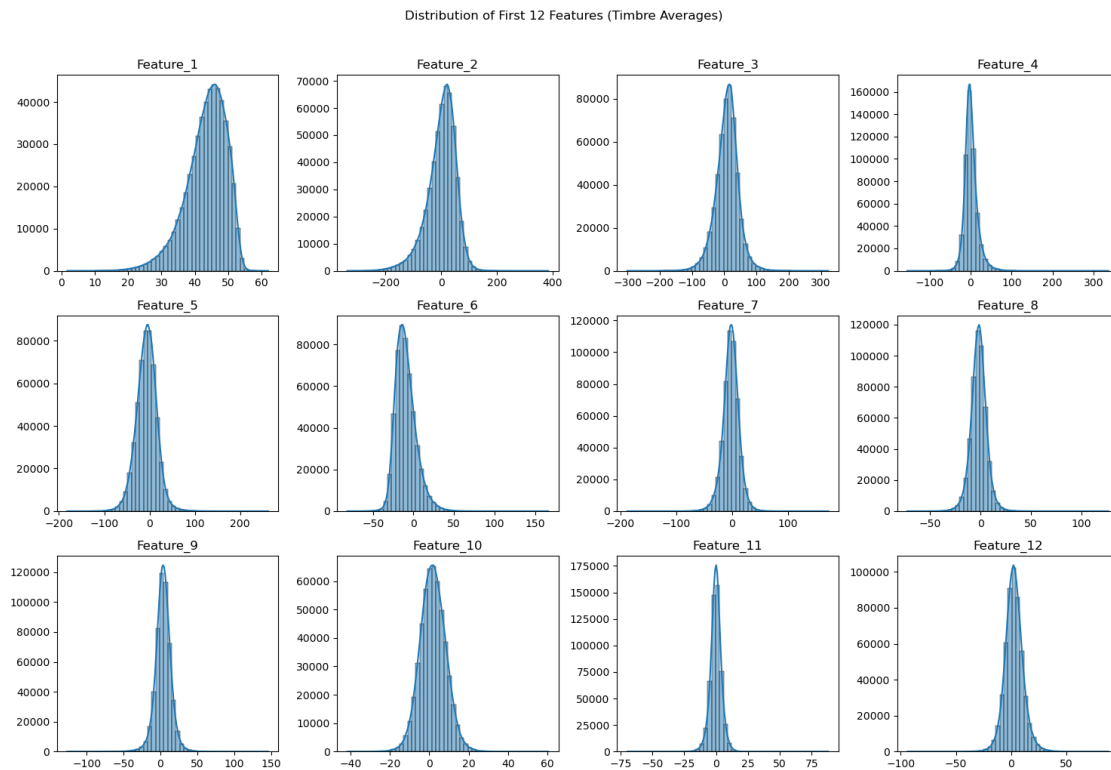
Observations: The dataset is heavily imbalanced, with a vast majority of songs from the 2000s, followed by the 1990s. Earlier decades have significantly fewer samples.

```
[6]: # --- 2. Feature Distribution Analysis ---
print("\n--- 2. Feature Distribution Analysis ---")
# Select a subset of features for detailed analysis (e.g., first 12, often
#   ↳ timbre averages)
# and maybe a few from the covariance features later on.
features_to_plot = df_raw.columns[1:13] # Features 1 to 12 (Timbre Averages)
print(f"Plotting distributions for features: {list(features_to_plot)}")
```

--- 2. Feature Distribution Analysis ---

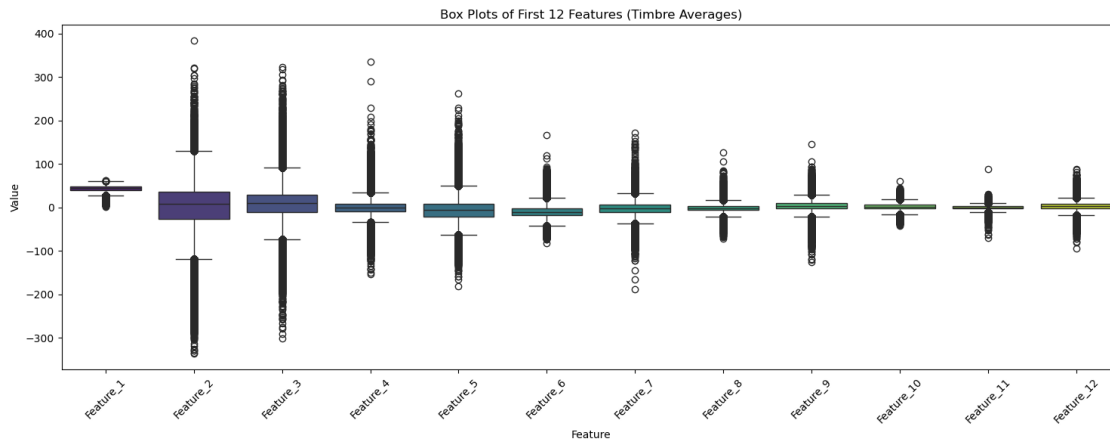
Plotting distributions for features: ['Feature_1', 'Feature_2', 'Feature_3', 'Feature_4', 'Feature_5', 'Feature_6', 'Feature_7', 'Feature_8', 'Feature_9', 'Feature_10', 'Feature_11', 'Feature_12']

```
[7]: plt.figure(figsize=(15, 10))
for i, col in enumerate(features_to_plot):
    plt.subplot(3, 4, i + 1) # Adjust grid size (3x4) as needed
    sns.histplot(df_raw[col], kde=True, bins=50)
    plt.title(col)
    plt.xlabel('')
    plt.ylabel('')
plt.suptitle('Distribution of First 12 Features (Timbre Averages)', y=1.02)
plt.tight_layout()
plt.savefig(os.path.join(PLOT_SAVE_DIR, 'eda_feature_distributions_hist.png'))
plt.show()
```



```
[8]: # Box plots can also show distribution and outliers
plt.figure(figsize=(15, 6))
sns.boxplot(data=df_raw[features_to_plot], palette='viridis')
plt.title('Box Plots of First 12 Features (Timbre Averages)')
plt.xlabel('Feature')
```

```
plt.ylabel('Value')
plt.xticks(rotation=45)
plt.tight_layout()
plt.savefig(os.path.join(PLOT_SAVE_DIR, 'eda_feature_distributions_box.png'))
plt.show()
```



```
[9]: print(f"\nObservations: Examine the plots for skewness, modality (number of
      ↳ peaks), and spread. Many features might appear roughly normally distributed,
      ↳ but could have long tails (indicating outliers).")
      print("Numerical summary:")
      print(df_raw[features_to_plot].describe())
```

Observations: Examine the plots for skewness, modality (number of peaks), and spread. Many features might appear roughly normally distributed but could have long tails (indicating outliers).

Numerical summary:

	Feature_1	Feature_2	Feature_3	Feature_4	\
count	515345.000000	515345.000000	515345.000000	515345.000000	
mean	43.387126	1.289554	8.658347	1.164124	
std	6.067558	51.580351	35.268585	16.322790	
min	1.749000	-337.092500	-301.005060	-154.183580	
25%	39.954690	-26.059520	-11.462710	-8.487500	
50%	44.258500	8.417850	10.476320	-0.652840	
75%	47.833890	36.124010	29.764820	8.787540	
max	61.970140	384.065730	322.851430	335.771820	

	Feature_5	Feature_6	Feature_7	Feature_8	\
count	515345.000000	515345.000000	515345.000000	515345.000000	
mean	-6.553601	-9.521975	-2.391089	-1.793236	
std	22.860785	12.857751	14.571873	7.963827	
min	-181.953370	-81.794290	-188.214000	-72.503850	

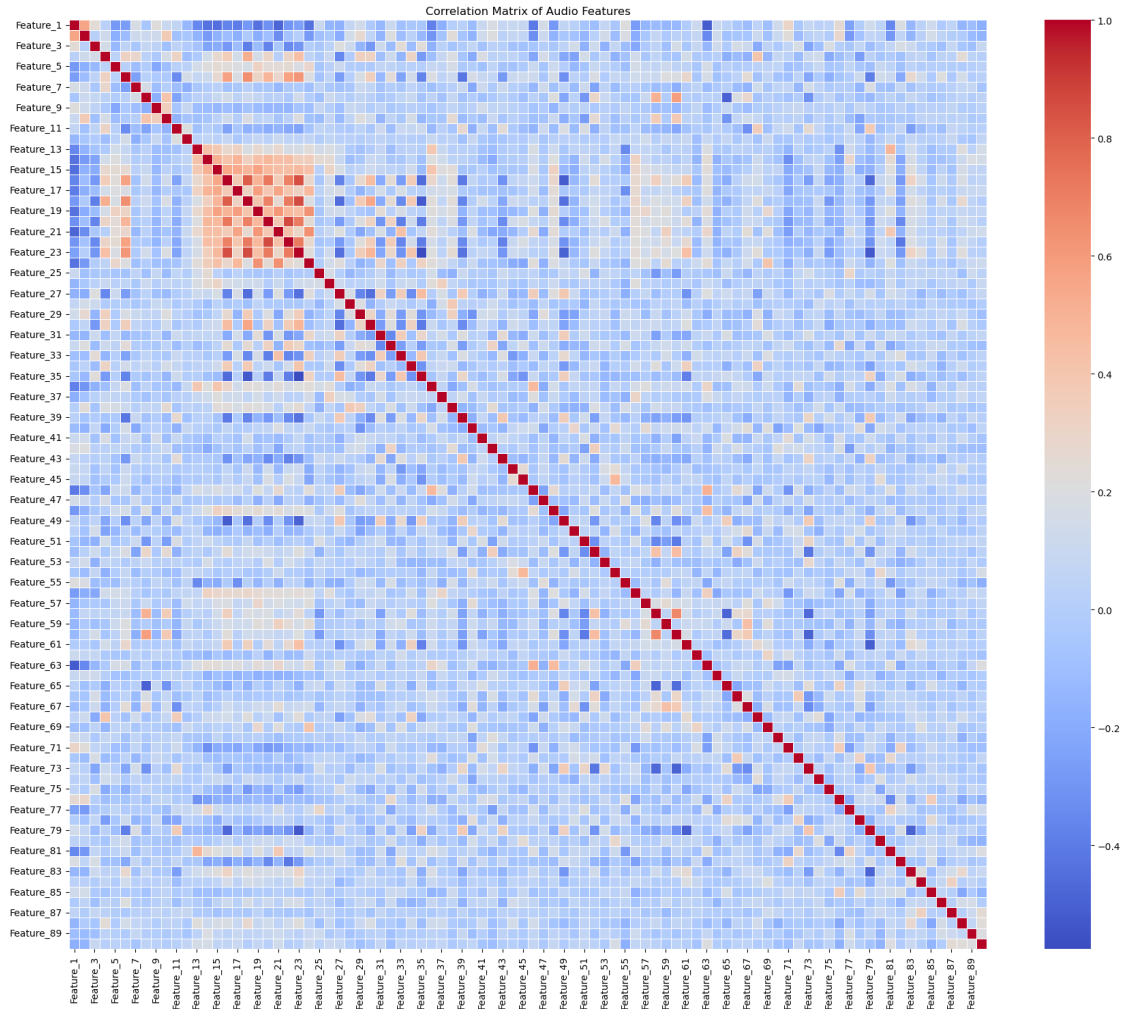
25%	-20.666450	-18.440990	-10.780600	-6.468420
50%	-6.007770	-11.188390	-2.046670	-1.736450
75%	7.741870	-2.388960	6.508580	2.913450
max	262.068870	166.236890	172.402680	126.741270

	Feature_9	Feature_10	Feature_11	Feature_12
count	515345.000000	515345.000000	515345.000000	515345.000000
mean	3.727876	1.882385	-0.146527	2.546063
std	10.582861	6.530232	4.370848	8.320190
min	-126.479040	-41.631660	-69.680870	-94.041960
25%	-2.293660	-2.444850	-2.652090	-2.550060
50%	3.822310	1.783520	-0.097950	2.313700
75%	9.961820	6.147220	2.435660	7.360330
max	146.297950	60.345350	88.020820	87.913240

```
[10]: # --- 3. Correlation Analysis ---
print("\n--- 3. Correlation Analysis ---")
# Calculate correlation matrix for numerical features (excluding Year and
↳ derived decade cols)
feature_cols = [col for col in df_raw.columns if col.startswith('Feature_')]
correlation_matrix = df_raw[feature_cols].corr()

plt.figure(figsize=(18, 15))
sns.heatmap(correlation_matrix, cmap='coolwarm', annot=False, fmt=".1f",
↳ linewidths=.5) # annot=True is too crowded for 90 features
plt.title('Correlation Matrix of Audio Features')
plt.tight_layout()
plt.savefig(os.path.join(PLOT_SAVE_DIR, 'eda_correlation_heatmap.png'))
plt.show()
```

--- 3. Correlation Analysis ---



```
[13]: # Find highly correlated pairs (optional)
threshold = 0.8
highly_correlated = correlation_matrix[abs(correlation_matrix) > threshold]
# Stack to get pairs, remove self-correlation, drop duplicates
corr_pairs = highly_correlated.unstack().sort_values(ascending=False).
    ↪ drop_duplicates()
corr_pairs = corr_pairs[corr_pairs != 1.0] # Remove self-correlations

print(f"\nHighly Correlated Feature Pairs (Threshold > {threshold}):")
if not corr_pairs.empty:
    print(corr_pairs)
else:
    print("No feature pairs found with correlation above the threshold.")
```

```
print(f"\nObservations: Look for blocks of high correlation (positive or
↪negative) in the heatmap. High correlation might suggest redundancy, but NNs
↪can sometimes handle it. The first 12 features (timbre averages) might show
↪some correlation among themselves, as might the covariance features.")
```

Highly Correlated Feature Pairs (Threshold > 0.8):

```
Feature_22  Feature_20    0.865684
Feature_18  Feature_23    0.859569
Feature_16  Feature_23    0.846649
           Feature_18    0.809554
Feature_1   Feature_2      NaN
dtype: float64
```

Observations: Look for blocks of high correlation (positive or negative) in the heatmap. High correlation might suggest redundancy, but NNs can sometimes handle it. The first 12 features (timbre averages) might show some correlation among themselves, as might the covariance features.

```
[14]: # --- 4. Outlier Analysis (using Box Plots from Feature Distribution) ---
print("\n--- 4. Outlier Analysis ---")
print("Refer back to the box plots generated in the 'Feature Distribution
↪Analysis' section.")
print("Box plots visually indicate potential outliers as points beyond the
↪'whiskers'.")
```

--- 4. Outlier Analysis ---

Refer back to the box plots generated in the 'Feature Distribution Analysis' section.

Box plots visually indicate potential outliers as points beyond the 'whiskers'.

```
[15]: # Example: Calculate IQR bounds for one feature
feature_example = 'Feature_1'
Q1 = df_raw[feature_example].quantile(0.25)
Q3 = df_raw[feature_example].quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
```

```
[16]: outliers = df_raw[(df_raw[feature_example] < lower_bound) |
↪(df_raw[feature_example] > upper_bound)]
print(f"\nExample Outlier Check for '{feature_example}':")
print(f"   IQR: {IQR:.2f}")
print(f"   Lower Bound (Q1 - 1.5*IQR): {lower_bound:.2f}")
print(f"   Upper Bound (Q3 + 1.5*IQR): {upper_bound:.2f}")
print(f"   Number of potential outliers (based on 1.5*IQR rule):
↪{len(outliers)}")
```

```
print(f" Percentage of potential outliers: {len(outliers) / len(df_raw) * 100:.  
↪2f}%")
```

Example Outlier Check for 'Feature_1':

IQR: 7.88

Lower Bound (Q1 - 1.5*IQR): 28.14

Upper Bound (Q3 + 1.5*IQR): 59.65

Number of potential outliers (based on 1.5*IQR rule): 10627

Percentage of potential outliers: 2.06%

```
[17]: print(f"\nObservations & Handling Strategy:")  
print(" - Many features show points beyond the 1.5*IQR whiskers, suggesting the  
↪presence of outliers.")  
print(" - Strategy Decision: For this project, we used StandardScaler in  
↪data_processing.py. While StandardScaler is sensitive to outliers, deep  
↪learning models (especially with techniques like Batch Norm, which we might  
↪test later) can sometimes be relatively robust.")  
print(" - Alternative strategies (not implemented here but considered):")  
print("   - Use RobustScaler: Scales using percentiles, less sensitive to  
↪outliers.")  
print("   - Clipping: Cap feature values at certain percentiles (e.g., 1st and  
↪99th).")  
print("   - Transformation: Apply log or Box-Cox transforms if features are  
↪highly skewed.")  
print(" - Chosen Approach: Proceed with StandardScaler, acknowledging the  
↪presence of outliers. We will monitor model performance and may revisit  
↪outlier handling if necessary.")
```

Observations & Handling Strategy:

- Many features show points beyond the 1.5*IQR whiskers, suggesting the presence of outliers.

- Strategy Decision: For this project, we used StandardScaler in data_processing.py. While StandardScaler is sensitive to outliers, deep learning models (especially with techniques like Batch Norm, which we might test later) can sometimes be relatively robust.

- Alternative strategies (not implemented here but considered):

- Use RobustScaler: Scales using percentiles, less sensitive to outliers.

- Clipping: Cap feature values at certain percentiles (e.g., 1st and 99th).

- Transformation: Apply log or Box-Cox transforms if features are highly skewed.

- Chosen Approach: Proceed with StandardScaler, acknowledging the presence of outliers. We will monitor model performance and may revisit outlier handling if necessary.

```
[18]: # --- 5. Categorical Features ---
print("\n--- 5. Categorical Features ---")
# Check data types again after loading
print(df_raw.info())
# Identify non-numeric columns (excluding our derived Decade_Name)
categorical_cols = df_raw.select_dtypes(include=['object', 'category']).columns
print(f"\nPotential categorical columns detected (excluding Decade_Name):_
↳{list(categorical_cols.drop('Decade_Name', errors='ignore'))}")
print("Observations: As expected for this dataset, all original predictor_
↳columns (Feature_1 to Feature_90) are numeric (float64). No categorical_
↳feature embedding strategy is required for the predictors.")
```

```
--- 5. Categorical Features ---
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 515345 entries, 0 to 515344
Data columns (total 94 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Year            515345 non-null  int64
1   Feature_1       515345 non-null  float64
2   Feature_2       515345 non-null  float64
3   Feature_3       515345 non-null  float64
4   Feature_4       515345 non-null  float64
5   Feature_5       515345 non-null  float64
6   Feature_6       515345 non-null  float64
7   Feature_7       515345 non-null  float64
8   Feature_8       515345 non-null  float64
9   Feature_9       515345 non-null  float64
10  Feature_10      515345 non-null  float64
11  Feature_11      515345 non-null  float64
12  Feature_12      515345 non-null  float64
13  Feature_13      515345 non-null  float64
14  Feature_14      515345 non-null  float64
15  Feature_15      515345 non-null  float64
16  Feature_16      515345 non-null  float64
17  Feature_17      515345 non-null  float64
18  Feature_18      515345 non-null  float64
19  Feature_19      515345 non-null  float64
20  Feature_20      515345 non-null  float64
21  Feature_21      515345 non-null  float64
22  Feature_22      515345 non-null  float64
23  Feature_23      515345 non-null  float64
24  Feature_24      515345 non-null  float64
25  Feature_25      515345 non-null  float64
26  Feature_26      515345 non-null  float64
27  Feature_27      515345 non-null  float64
28  Feature_28      515345 non-null  float64
```

29	Feature_29	515345	non-null	float64
30	Feature_30	515345	non-null	float64
31	Feature_31	515345	non-null	float64
32	Feature_32	515345	non-null	float64
33	Feature_33	515345	non-null	float64
34	Feature_34	515345	non-null	float64
35	Feature_35	515345	non-null	float64
36	Feature_36	515345	non-null	float64
37	Feature_37	515345	non-null	float64
38	Feature_38	515345	non-null	float64
39	Feature_39	515345	non-null	float64
40	Feature_40	515345	non-null	float64
41	Feature_41	515345	non-null	float64
42	Feature_42	515345	non-null	float64
43	Feature_43	515345	non-null	float64
44	Feature_44	515345	non-null	float64
45	Feature_45	515345	non-null	float64
46	Feature_46	515345	non-null	float64
47	Feature_47	515345	non-null	float64
48	Feature_48	515345	non-null	float64
49	Feature_49	515345	non-null	float64
50	Feature_50	515345	non-null	float64
51	Feature_51	515345	non-null	float64
52	Feature_52	515345	non-null	float64
53	Feature_53	515345	non-null	float64
54	Feature_54	515345	non-null	float64
55	Feature_55	515345	non-null	float64
56	Feature_56	515345	non-null	float64
57	Feature_57	515345	non-null	float64
58	Feature_58	515345	non-null	float64
59	Feature_59	515345	non-null	float64
60	Feature_60	515345	non-null	float64
61	Feature_61	515345	non-null	float64
62	Feature_62	515345	non-null	float64
63	Feature_63	515345	non-null	float64
64	Feature_64	515345	non-null	float64
65	Feature_65	515345	non-null	float64
66	Feature_66	515345	non-null	float64
67	Feature_67	515345	non-null	float64
68	Feature_68	515345	non-null	float64
69	Feature_69	515345	non-null	float64
70	Feature_70	515345	non-null	float64
71	Feature_71	515345	non-null	float64
72	Feature_72	515345	non-null	float64
73	Feature_73	515345	non-null	float64
74	Feature_74	515345	non-null	float64
75	Feature_75	515345	non-null	float64
76	Feature_76	515345	non-null	float64

```

77 Feature_77      515345 non-null float64
78 Feature_78      515345 non-null float64
79 Feature_79      515345 non-null float64
80 Feature_80      515345 non-null float64
81 Feature_81      515345 non-null float64
82 Feature_82      515345 non-null float64
83 Feature_83      515345 non-null float64
84 Feature_84      515345 non-null float64
85 Feature_85      515345 non-null float64
86 Feature_86      515345 non-null float64
87 Feature_87      515345 non-null float64
88 Feature_88      515345 non-null float64
89 Feature_89      515345 non-null float64
90 Feature_90      515345 non-null float64
91 Decade_Start    515345 non-null int64
92 Decade_Label    515345 non-null int64
93 Decade_Name     515345 non-null object
dtypes: float64(90), int64(3), object(1)
memory usage: 369.6+ MB
None

```

Potential categorical columns detected (excluding Decade_Name): []
Observations: As expected for this dataset, all original predictor columns (Feature_1 to Feature_90) are numeric (float64). No categorical feature embedding strategy is required for the predictors.

```

[19]: # --- Summary of EDA Findings ---
print("\n--- Summary of Key EDA Findings ---")
print("1.  **Target Variable (Decade):** Heavily imbalanced, dominated by 2000s
    and 1990s.")
print("2.  **Features:** All 90 predictor features are numeric (float).")
print("3.  **Distributions:** Feature distributions vary. Some are roughly
    normal, others might be skewed or have multiple peaks (visual inspection
    needed per feature).")
print("4.  **Correlations:** Some correlations exist between features,
    particularly noted visually within blocks (e.g., early timbre features,
    later covariance features). No extremely high correlations (>0.95) jumped
    out immediately in the sample check, but moderate correlations are present.")
print("5.  **Outliers:** Potential outliers detected in many features based on
    visual inspection of box plots and IQR checks.")
print("6.  **Missing Values:** No missing values detected by `df.info()`
    (consistent with dataset description).")
print("7.  **Preprocessing Decisions (Recap):**")
print("    - Decade binning successfully converted regression to classification.
    ")
print("    - Stratified splitting addressed the class imbalance during data
    partitioning.")

```

```
print("    - StandardScaler was used for feature scaling, acknowledging outlier_
presence.")
print("    - No categorical encoding needed for predictors.")
```

--- Summary of Key EDA Findings ---

1. **Target Variable (Decade):** Heavily imbalanced, dominated by 2000s and 1990s.
2. **Features:** All 90 predictor features are numeric (float).
3. **Distributions:** Feature distributions vary. Some are roughly normal, others might be skewed or have multiple peaks (visual inspection needed per feature).
4. **Correlations:** Some correlations exist between features, particularly noted visually within blocks (e.g., early timbre features, later covariance features). No extremely high correlations (>0.95) jumped out immediately in the sample check, but moderate correlations are present.
5. **Outliers:** Potential outliers detected in many features based on visual inspection of box plots and IQR checks.
6. **Missing Values:** No missing values detected by `df.info()` (consistent with dataset description).
7. **Preprocessing Decisions (Recap):**
 - Decade binning successfully converted regression to classification.
 - Stratified splitting addressed the class imbalance during data partitioning.
 - StandardScaler was used for feature scaling, acknowledging outlier presence.
 - No categorical encoding needed for predictors.